



Early Exit Probing on the 10 MMEB evaluation sets: deck vs. our reproduction

Native candidate sets. Ours never trains on classification data; thresholds are calibrated on held-out rows of each run's own training pool. A-OKVQA is scored on its 1,145-row val split.

Deck (5-dataset training pool) P1: ours, A-OKVQA only (500 steps) P3: ours, A-OKVQA + ScienceQA (1000 steps)



mean exit depth: deck | P1 | P3