


Qwen3-VL 2B & 8B on A-OKVQA: estimated training compute vs. accuracy
every non-linear-probe method family in logs/ (depth pruning x 3 patterns, FFN-head-only, LoRA+FFN-head, adapter-freeze sweep, single-seam ablations)



Note: compute is an estimate (this repo has no surviving per-run FLOP/wall-clock counters), not a measurement -- see docstring for the exact formula and its assumptions. All 44 runs use identical steps (500) and batch size (64), confirmed across every sweep driver log, so the proxy is directly proportional to relative training compute up to one shared constant. Metric is the in-batch 4-way accuracy proxy used throughout this repo's other figures (chance=25%), not the harder full-pool P@1 retrieval metric (chance~0.11%, full 2B model ~21%) computed separately in some p1_*.log files. Excludes eval_probe/ (linear probe, not a fine-tuning run) and collapse_pretraining_check/ (an untrained/random baseline, ~chance with a negative discriminative gap) per the request scope.