


Qwen3-VL 2B & 8B on A-OKVQA: estimated training compute vs. accuracy,
with the linear-probe floor added



Note: compute is an estimate, not a measurement (see docstring). All gradient-trained runs use identical steps (500) and batch size (64). The linear probe (teal) is NOT one of them -- it's a single forward pass over the training set plus a closed-form/few-iteration linear fit, zero backward passes through the backbone, so its plotted x is a conservative (upper-bound) estimate; its true compute cost relative to everything else is almost certainly even lower than shown. The probe's accuracy is also a different metric -- real question_accuracy vs. the in-batch 4-way proxy every other point uses (both chance=25%, but not the same measurement) -- so read its vertical position as indicative, not a precise apples-to-apples comparison to the other curves. Excludes eval_probe/'s other configs (this is the keep-first-N probe sweep only) and collapse_pretraining_check/ (an untrained/random baseline).