



Qwen3-VL-2B, full unpruned depth (28 layers): LoRA vs. FFN-head-only

Train loss: full-depth 2B, LoRA vs. FFN-head-only



Full-depth accuracy



Note: LoRA's best/latest (84.0/85.0%) are read from pruning_accuracy_v2.png's own printed labels at x=0% pruned (no raw log survived for this point); everything else is exact from this session's logs. Same crossover trend as the N=3-9 sweep (plot_2b_ffnhead_vs_lora.png) but wider: LoRA's lead over FFN-head-only grows with more depth.