



Qwen3-VL-2B, full unpruned depth (28 layers): LoRA vs. FFN-head vs. combined

Train loss: full-depth 2B, all three methods



Full-depth accuracy



Note: LoRA-only's best/latest (84.0/85.0%) are read from pruning_accuracy_v2.png's own printed labels at x=0% pruned (no raw log survived for this point); everything else is exact from this session's logs. Combining LoRA + FFN-head beats LoRA alone on best-checkpoint accuracy (+4.0pt) but shows more checkpoint-to-checkpoint variance (88.0->84.5->85.0 across steps 380-440) and lands roughly even with LoRA-only at the final step (84.5% vs 85.0%).