


Qwen3-VL-2B depth pruning vs A-OKVQA accuracy

alternating drop (cliff) vs. contiguous keep-first-N vs. the real linear-probe floor



Note: all curves except "linear probe" use the in-batch 4-way training-accuracy proxy metric; the linear-probe curve uses real A-OKVQA question_accuracy (a properly-evaluated multiple-choice metric). Both are chance=25%, but treat cross-curve deltas as indicative, not exactly equivalent (see compute_train_accuracy.py vs. the real eval).